

```

# For data operations
import pandas as pd

# For visualisations
import seaborn as sns

# For visualisations
import matplotlib.pyplot as plt

# To set the styles for the visualisations
from matplotlib import style
style.use('seaborn')

# To find Optimal parameters for hyperparameter tuning
from sklearn.model_selection import GridSearchCV

# To split the data into training and testing sets
from sklearn.model_selection import train_test_split

# To scale the Data
from sklearn.preprocessing import StandardScaler

# importing various Algorithms
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from xgboost import XGBRegressor
from sklearn.neighbors import KNeighborsClassifier

# To measure the model performance
from sklearn.metrics import mean_absolute_error, r2_score
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
import warnings
warnings.filterwarnings('ignore')

# Reading the data from the csv file
car_data = pd.read_csv('C:/Users/Shubham/OneDrive - Teesside
University/Machine Learning Lab/ford.csv')

# Getting the idea of the data we are dealing with
car_data.head()

```

	model	year	price	transmission	mileage	fuelType	tax	mpg
engineSize								
0	Fiesta	2017	12000	Automatic	15944	Petrol	150	57.7
1.0								
1	Focus	2018	14000	Manual	9083	Petrol	150	57.7
1.0								
2	Focus	2017	13000	Manual	12456	Petrol	150	57.7

```

1.0
3   Fiesta   2019   17500      Manual   10460   Petrol   145   40.3
1.5
4   Fiesta   2019   16500   Automatic   1482    Petrol   145   48.7
1.0

```

```
car_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17966 entries, 0 to 17965
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   model            17966 non-null  object
1   year             17966 non-null  int64
2   price            17966 non-null  int64
3   transmission     17966 non-null  object
4   mileage          17966 non-null  int64
5   fuelType         17966 non-null  object
6   tax              17966 non-null  int64
7   mpg              17966 non-null  float64
8   engineSize       17966 non-null  float64
dtypes: float64(2), int64(4), object(3)
memory usage: 1.2+ MB

```

```

# checking any null entries
car_data.isnull().sum()

```

```

model            0
year             0
price            0
transmission     0
mileage          0
fuelType         0
tax              0
mpg              0
engineSize       0
dtype: int64

```

```

# Getting statistical information about the dataset
car_data.describe()

```

	year	price	mileage	tax
mpg \				
count	17966.000000	17966.000000	17966.000000	17966.000000
mean	2016.866470	12279.534844	23362.608761	113.329456
std	2.050336	4741.343657	19472.054349	62.012456
min	1996.000000	495.000000	1.000000	0.000000

```

20.800000
25%      2016.000000    8999.000000    9987.000000    30.000000
52.300000
50%      2017.000000   11291.000000   18242.500000   145.000000
58.900000
75%      2018.000000   15299.000000   31060.000000   145.000000
65.700000
max       2060.000000   54995.000000  177644.000000   580.000000
201.800000

```

```

engineSize
count    17966.000000
mean       1.350807
std        0.432367
min        0.000000
25%        1.000000
50%        1.200000
75%        1.500000
max        5.000000

```

```
car_data.columns
```

```

Index(['model', 'year', 'price', 'transmission', 'mileage',
       'fuelType', 'tax',
       'mpg', 'engineSize'],
      dtype='object')


```

```
# Exploring the categorical data of model types
```

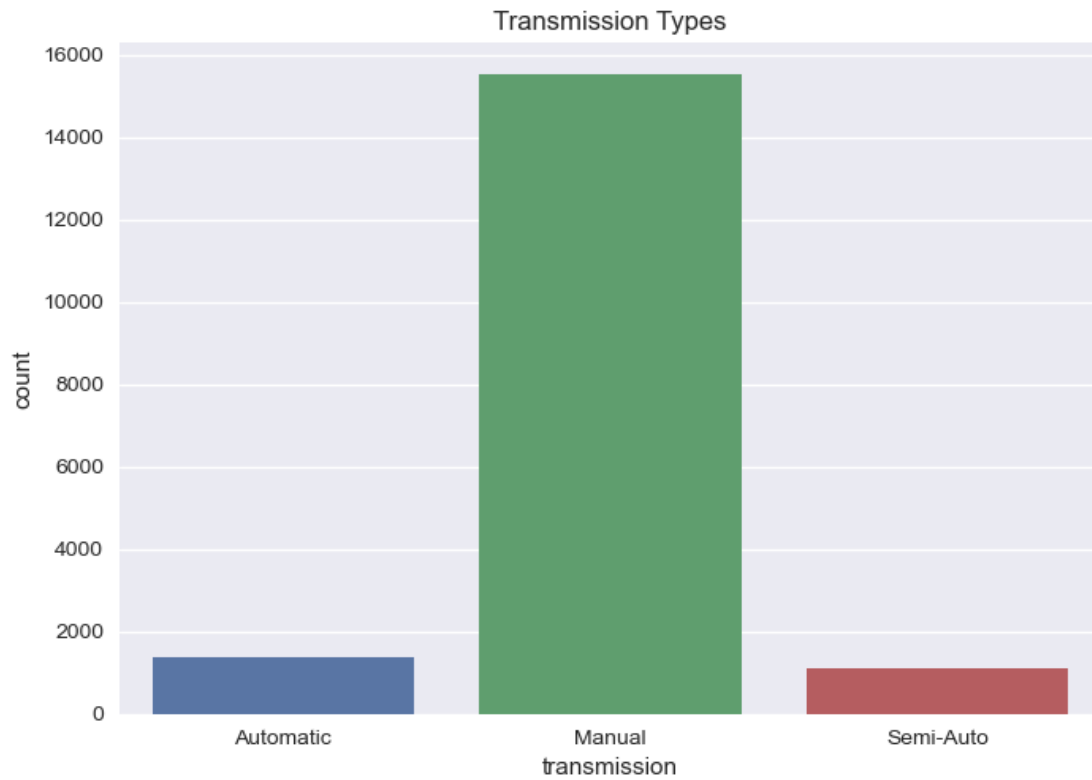
```

plt.figure(figsize=(12,10))
sns.countplot(y='model', data=car_data)
plt.title('Model types')
plt.show()

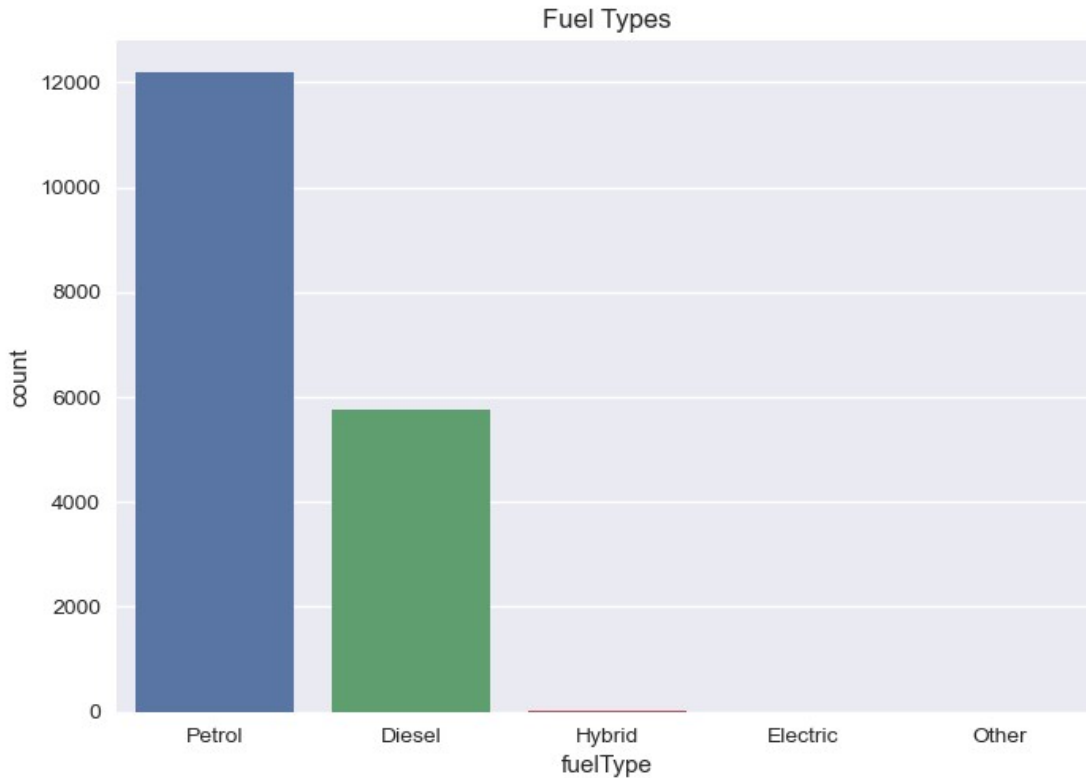
```



```
# Exploring the categorical data of transmission types
sns.countplot(x='transmission', data=car_data)
plt.title('Transmission Types')
plt.show()
```



```
# Exploring the categorical data of fuel types  
sns.countplot(x='fuelType', data=car_data)  
plt.title('Fuel Types')  
plt.show()
```



To see numerical values of different categorical data

```
print(car_data['model'].value_counts())
print("\n\n")
print(car_data['transmission'].value_counts())
print("\n\n")
print(car_data['fuelType'].value_counts())
```

Fiesta	6557
Focus	4588
Kuga	2225
EcoSport	1143
C-MAX	543
Ka+	531
Mondeo	526
B-MAX	355
S-MAX	296
Grand C-MAX	247
Galaxy	228
Edge	208
KA	199
Puma	80
Tourneo Custom	69
Grand Tourneo Connect	59
Mustang	57
Tourneo Connect	33
Fusion	16

```

Streetka          2
Ranger            1
Escort            1
Transit Tourneo   1
Focus             1
Name: model, dtype: int64

```

```

Manual          15518
Automatic        1361
Semi-Auto       1087
Name: transmission, dtype: int64

```

```

Petrol          12179
Diesel           5762
Hybrid           22
Electric          2
Other             1
Name: fuelType, dtype: int64


```

```

# Creating a variable to store the data from the column
fuelType = car_data['fuelType']
transmission = car_data['transmission']
price = car_data['price']
fig, axes = plt.subplots(1,2, figsize=(15,5), sharey=True)
fig.suptitle('Visualizing categorical data columns')
sns.barplot(x=fuelType, y=price, ax=axes[0])
sns.barplot(x=transmission, y=price, ax = axes[1])

<AxesSubplot:xlabel='transmission', ylabel='price'>

```



```


# performing ordinal encoding on the data to convert into numerical
values.

```

```
car_data.replace({'transmission':{'Manual':0, 'Automatic':1, 'Semi-
Auto':2}}, inplace=True)
car_data.replace({'fuelType':{'Petrol':0, 'Diesel':1, 'Hybrid':2,
'Electric':3, 'Other':4}}, inplace=True)
```

car price value based on model


```
ax = sns.boxplot(data=car_data, x=car_data['model'],
y=car_data['price'])
ax.tick_params(axis='x', rotation=90)
```




car price value based on mileage driven

```
sns.jointplot(data=car_data, x=car_data['mileage'],
y=car_data['price'])
```

<seaborn.axisgrid.JointGrid at 0x1ef52af52b0>



```
# car price value based on manufacured year  
ax = sns.barplot(data=car_data, x=car_data['year'],  
y=car_data['price'])  
ax.tick_params(axis='x', rotation=90)
```



```
# pairplot
sns.pairplot(car_data)

<seaborn.axisgrid.PairGrid at 0x1ef520a0550>
```



using the head function to see the modified data

```
car_data.head()
```

dropping the model column

```
car_data = car_data.drop("model", axis=1)
```


visualising heat map to show correlation between attributes and the target column

```
plt.figure(figsize=(10,7))
```

```
sns.heatmap(car_data.corr(), annot=True)
```

```
plt.title('Correlation between the columns')
```

```
plt.show()
```



sorting the correlation of the price column

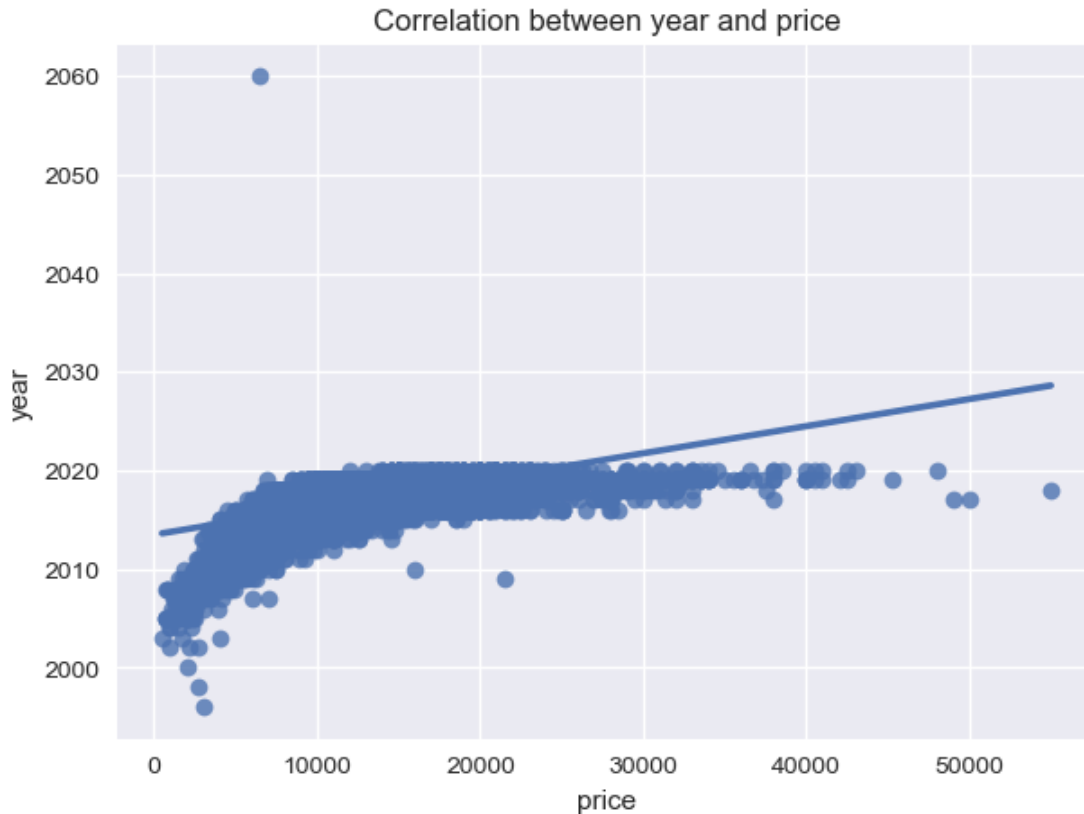
```
car_data.corr()['price'].sort_values()
```

```
mileage      -0.530659
mpg          -0.346419
fuelType      0.209225
transmission  0.231806
tax           0.406857
engineSize    0.411178
year          0.636009
price         1.000000
Name: price, dtype: float64
```

plotting regression plot

```
fig = plt.figure(figsize=(7,5))
plt.title('Correlation between year and price')
sns.regplot(x='price', y='year', data=car_data)
```

```
<AxesSubplot:title={'center': 'Correlation between year and price'},
xlabel='price', ylabel='year'>
```



```
# splitting the data into x and y
X = car_data.drop('price', axis=1)
y = car_data['price']

# printing the shape of x and y
print("Shape of X is :", X.shape)
print("Shape of y is :", y.shape)

Shape of X is : (17966, 7)
Shape of y is : (17966,)

# splitting the data into training and testing sets.
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# printing the shape of training and testing data
print("Shape of X_train is: ", X_train.shape)
print("Shape of y_train is: ", y_train.shape)
print("Shape of X_test is: ", X_test.shape)
print("Shape of y_test is: ", y_test.shape)

Shape of X_train is: (14372, 7)
Shape of y_train is: (14372,)
Shape of X_test is: (3594, 7)
Shape of y_test is: (3594,)
```

```

# normalising the data
scaler = StandardScaler()

# applying normalisation on X data
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# creating the linear regression model
linreg = LinearRegression()
linreg.fit(X_train, y_train)
linreg_pred = linreg.predict(X_test)

# evaluating the linear regression model performance
linreg_mae = mean_absolute_error(y_test, linreg_pred)
linreg_r2 = r2_score(y_test, linreg_pred)
print("MAE of linear regression model is:", linreg_mae)
print("R2 score of linear regression model is:", linreg_r2)

MAE of linear regression model is: 1778.2319322410349
R2 score of linear regression model is: 0.7379425731911509

# performing cross validation on test data
linreg_score = cross_val_score(linreg, X_test, y_test, cv=4)
print("Linear Regression model accuracy is:
{}".format(linreg_score.mean()*100))

Linear Regression model accuracy is: 73.87545898136361

# Defining the hyperparameters to tune
parameters = {'fit_intercept': [True, False],
              'normalize': [True, False],
              'copy_X': [True, False]}

# Creating the GridSearchCV object for linear regression model
linreg_cv = GridSearchCV(LinearRegression(), param_grid=parameters,
cv=5)

# Fitting the GridSearchCV object to training data
linreg_cv.fit(X_train, y_train)

GridSearchCV(cv=5, estimator=LinearRegression(),
              param_grid={'copy_X': [True, False],
                          'fit_intercept': [True, False],
                          'normalize': [True, False]})

# Printing the best hyperparameters found
print("Best hyperparameters:", linreg_cv.best_params_)

Best hyperparameters: {'copy_X': True, 'fit_intercept': True,
'normalize': True}

# Using the best hyperparameters to predict on the test data
linreg_pred = linreg_cv.predict(X_test)

```

```
# Calculating the MAE and R2 score for the model with the best hyperparameters
linreg_mae = mean_absolute_error(y_test, linreg_pred)
linreg_r2 = r2_score(y_test, linreg_pred)
print("MAE of linear regression model with best hyperparameters:",
linreg_mae)
print("R2 score of linear regression model with best hyperparameters:", linreg_r2)
```

MAE of linear regression model with best hyperparameters:
1778.2319322410349
R2 score of linear regression model with best hyperparameters:
0.7379425731911511

```
# creating the decision model
dtree = DecisionTreeRegressor()
dtree.fit(X_train, y_train)
dtree_pred = dtree.predict(X_test)
```

```
# evaluating the decision tree model performance
dtree_mae = mean_absolute_error(y_test, dtree_pred)
dtree_r2 = r2_score(y_test, dtree_pred)
print("MAE of decision tree model is:", dtree_mae)
print("R2 score of decision tree model is:", dtree_r2)
```

MAE of decision tree model is: 1174.6750602856614
R2 score of decision tree model is: 0.8597171999082486

```
# performing cross validation on test data
dtree_score = cross_val_score(dtree, X_test, y_test, cv=4)
print("Decision Tree model accuracy is:
{}".format(dtree_score.mean()*100))
```

Decision Tree model accuracy is: 82.3626463016493

```
# Defining the hyperparameters for the decision tree model
parameters = {
    'max_depth': [3, 5, 7, 10],
    'min_samples_split': [2, 5, 10, 15],
    'min_samples_leaf': [1, 2, 5, 10]
}
```

```
# Creating the GridSearchCV object for decision tree model
dtree_cv = GridSearchCV(DecisionTreeRegressor(),
param_grid=parameters, cv=5)
```

```
# Fitting the GridSearchCV object to training data
dtree_cv.fit(X_train, y_train)
```

```
GridSearchCV(cv=5, estimator=DecisionTreeRegressor(),
              param_grid={'max_depth': [3, 5, 7, 10],
```

```

        'min_samples_leaf': [1, 2, 5, 10],
        'min_samples_split': [2, 5, 10, 15]})

# Printing the best hyperparameters found
print("Best hyperparameters:", dtree_cv.best_params_)

Best hyperparameters: {'max_depth': 10, 'min_samples_leaf': 5,
'min_samples_split': 15}

# Using the best hyperparameters to predict on the test data
dtree_pred = dtree_cv.predict(X_test)

# Calculating the MAE and R2 score for the model with the best
hyperparameters
dtree_mae = mean_absolute_error(y_test, dtree_pred)
dtree_r2 = r2_score(y_test, dtree_pred)
print("MAE of decision tree model with best hyperparameters:",
dtree_mae)
print("R2 score of decision tree model with best hyperparameters:",
dtree_r2)

MAE of decision tree model with best hyperparameters:
1007.3313843147844
R2 score of decision tree model with best hyperparameters:
0.9031300534856432

# creating the XGBoost model
xgb = XGBRegressor()
xgb.fit(X_train, y_train)
xgb_prediction = xgb.predict(X_test)

# evaluating the XGBoost model performance
xgb_mae = mean_absolute_error(y_test, xgb_prediction)
xgb_r2 = r2_score(y_test, xgb_prediction)
print("MAE of xgboost model is:", xgb_mae)
print("R2 score of xgboost model is:", xgb_r2)

MAE of xgboost model is: 914.1581133384472
R2 score of xgboost model is: 0.9196627631595996

# performing cross validation on test data
xgb_score = cross_val_score(xgb, X_test, y_test, cv=4)
print("xgboost model accuracy is: {}".format(xgb_score.mean()*100))

xgboost model accuracy is: 89.60526998355522

#printing best cross validation score
print("best cross validation score: {:.2f}".format(grid.best_score_))
print("best parameters: ", grid.best_params_)

```

NameError

Traceback (most recent call


```

last)
~\AppData\Local\Temp\ipykernel_20908\3877909406.py in <module>
      1 #printing best cross validation score
----> 2 print("best cross validation score:
{:.2f}".format(grid.best_score_))
      3 print("best parameters: ", grid.best_params_)

NameError: name 'grid' is not defined

# Defining the hyperparameters to tune
param_grid = {
    'learning_rate': [0.1, 0.01, 0.001],
    'max_depth': [3, 4, 5],
    'n_estimators': [50, 100, 150]
}

# Creating a GridSearchCV object with the XGBoost model and the
hyperparameters
xgb_grid = GridSearchCV(XGBRegressor(), param_grid=param_grid, cv=5)

# Fitting the GridSearchCV object to the training data
xgb_grid.fit(X_train, y_train)

GridSearchCV(cv=5,
              estimator=XGBRegressor(base_score=None, booster=None,
                                     callbacks=None,
                                     colsample_bylevel=None,
                                     colsample_bynode=None,
                                     colsample_bytree=None,
                                     early_stopping_rounds=None,
                                     enable_categorical=False,
                                     eval_metric=None,
                                     feature_types=None,
                                     gamma=None,
                                     gpu_id=None,
                                     grow_policy=None,
                                     importance_type=None,
                                     interaction_constraints=None,
                                     learning_rate=None,
                                     max_bin=None,
                                     max_cat_threshold=None,
                                     max_cat_to_onehot=None,
                                     max_delta_step=None,
                                     max_depth=None,
                                     max_leaves=None,
                                     min_child_weight=None,
                                     missing=nan,
                                     monotone_constraints=None,
                                     n_estimators=100,
                                     n_jobs=None,
                                     num_parallel_tree=None,
                                     predictor=None,
                                     random_state=None,
                                     ...),
              param_grid={'learning_rate': [0.1, 0.01, 0.001],

```

```

        'max_depth': [3, 4, 5],
        'n_estimators': [50, 100, 150]})

# Printing the best hyperparameters found
print("Best hyperparameters:", xgb_grid.best_params_)

Best hyperparameters: {'learning_rate': 0.1, 'max_depth': 5,
'n_estimators': 150}

# Using the best hyperparameters to predict on the test data
xgb_pred = xgb_grid.predict(X_test)

# Calculating the MAE and R2 score for the model with the best
hyperparameters
xgb_mae = mean_absolute_error(y_test, xgb_pred)
xgb_r2 = r2_score(y_test, xgb_pred)
print("MAE of XGBoost model with best hyperparameters:", xgb_mae)
print("R2 score of XGBoost model with best hyperparameters:", xgb_r2)

MAE of XGBoost model with best hyperparameters: 932.5599693506375
R2 score of XGBoost model with best hyperparameters:
0.9168357941605024

car_data.columns

Index(['year', 'price', 'transmission', 'mileage', 'fuelType', 'tax',
'mpg',
      'engineSize'],
      dtype='object')

car_data.head()

```

	year	price	transmission	mileage	fuelType	tax	mpg	engineSize
0	2017	12000	1	15944	0	150	57.7	1.0
1	2018	14000	0	9083	0	150	57.7	1.0
2	2017	13000	0	12456	0	150	57.7	1.0
3	2019	17500	0	10460	0	145	40.3	1.5
4	2019	16500	1	1482	0	145	48.7	1.0

```

# creating new data for the model to predict
data = {'year':2017, 'transmission':1, 'mileage':15944, 'fuelType':0,
'tax':150, 'mpg':57.7,
      'engineSize':1.0}
index= [0]
new_car_data = pd.DataFrame(data, index)
new_car_data

   year  transmission  mileage  fuelType  tax  mpg  engineSize
0  2017             1   15944         0  150  57.7          1.0

# implementing the data created to the XGB model
new_prediction = xgb_grid.predict(new_car_data)
print("The car price for the new data is: ", new_prediction)

```

The car price for the new data is: [11068.842]